

# **Solving Large-Scale Dynamic Vehicle Routing Problems with Stochastic Requests**

Alexandre Florio

VeRoLog Webinar, March 31, 2022



# Solving Large-Scale Dynamic Vehicle Routing with Stochastic Requests in Real-Time

Joint work with Jian Zhang, Kelin Luo and Tom Van Woensel

<https://arxiv.org/abs/2202.12983>

<https://github.com/amflorio/dvrp-stochastic-requests>

<https://youtu.be/D57xNfU73as>

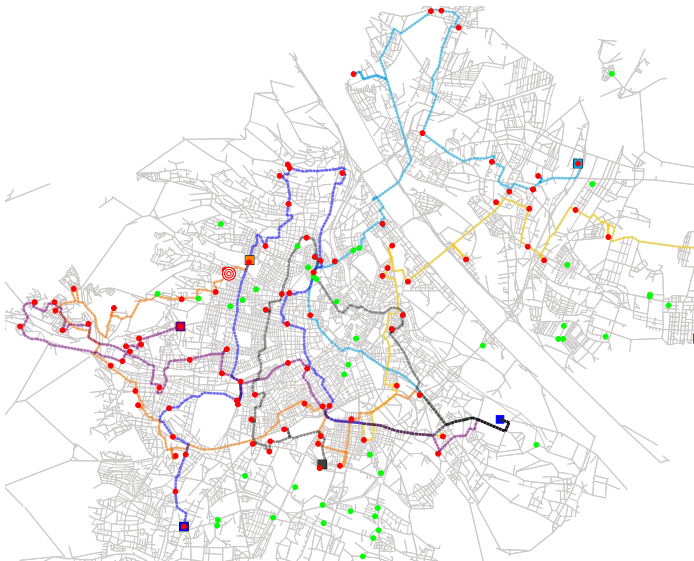


## General Setting & Motivation

- Dynamic VRP with stochastic requests (DVRPSR)
  - Static (planned) and dynamic (on-demand) requests
  - Applications: technician routing, package collection (first-mile, courier or marketplace sellers), field service management
- **Request:** pickup (or service) location and service time
  - Requests may be rejected
- Multiple drivers/vehicles
- Vehicles must return to the depot by the given deadline
- **Goal:** serve as many requests as possible



# General Setting & Motivation





## Related Literature

Literature	Problem setting				Initial plan		Online policy	
	RT	Req.	Veh.	Graph	Method	Ant.	Method	Ant.
Bent & Van Hentenryck (2004)	×	100	n/a	Synth.	MSA	✓	MSA	✓
Chen & Xu (2006)	×	100	n/a	Synth.	CGBH	×	CGBH	×
Hvattum et al. (2006)	×	130	n/a	Synth.	DSHH	✓	DSHH	✓
Ichoua et al. (2006)	✓	240	6	Synth.	TS	✓	TS	✓
Thomas (2007)	×	50	1	Synth.	GRASP	×	Heuristics	✓
Azi et al. (2012)	✓	144	5	Synth.	n/a		ALNS	✓
Ferrucci & Bock (2015, 2016)	×	150	12	SSSN	TS	✓	TS	✓
Klapp et al. (2018a)	×	40	1	Synth.	B&C	✓	Rollout	✓
Ulmer et al. (2018a)	×	100	1	Synth.	CI	×	VFA	✓
Ulmer et al. (2018b)	×	100	1	Synth.	CI	×	VFA	✓
Ulmer et al. (2019)	×	100	1	Synth.	CI	×	VFA+rollout	✓
van Heeswijk et al. (2019)	×	400	n/a	Synth., SSSN	CW	×	VFA, rollout	✓
Voccia et al. (2019)	×	192	13	Synth.	n/a		MSA	✓
Ulmer (2020)	✓	180	3	Synth.	n/a		VFA	✓
Ulmer & Thomas (2020)	✓	50	1	Synth.	n/a		VFA	✓
<b>This paper</b>	✓	<b>947</b>	<b>20</b>	<b>LSSN</b>	<b>PbCGBH</b>	✓	<b>PbPs</b>	✓

(For a review: Soeffker, Ulmer, Mattfeld (2021, EJOR): 10.1016/j.ejor.2021.07.014)



- Problem formulation: sequential stochastic optimization model
  - Decisions must be taken “real-time”, each time a request arrives
- Key contribution: approximation of the reward-to-go (potential) by relaxed knapsack models
  - Stochastic lookahead policy with **zero** tunable parameters
- First decision: initial plan: column generation-based petal heuristic
- **Remaining decisions: potential-based scheduling policy**



## Notation

$\mathcal{G} = (\mathcal{V}, \mathcal{A})$	Street network graph
$\mathcal{V}$	Road junctions or intersections
$\mathcal{A}$	Road segments
$0 \in \mathcal{V}$	Depot
$t(i, j)$	Duration of the fastest path from node $i$ to node $j$
$[0, U]$	Service period
$K$	Fleet size



## Notation

$\mathcal{D}(u)$	(Ordered) set of dynamic requests up to instant $u \in [0, U]$
$(u, i, d)$	Request, where $u \in [0, U]$ is the arrival/release time, $i \in \mathcal{V}$ is the location and $d \in \mathbb{R}_{>0}$ is the service duration
$T \equiv  \mathcal{D}(U) $	R.v. indicating the total number of dynamic requests

We assume the spatiotemporal probability distribution of dynamic requests can be sampled from (e.g., data is available)

$\omega = \{r_1^\omega, \dots, r_{T_\omega}^\omega\}$  Sample path of the request arrival process





## MDP Model: Decision Epochs and States

- Decision epoch: moment when a decision must be made
- In our setting, we have:
  - Initial decision: setup initial routes to serve static/planned requests
  - Following decisions: each time a dynamic request arrives, decide whether to accept or reject the request and, if accept, how to serve it
- Hence,  $1 + T$  decision epochs, where  $T$  is unknown

### Definition (Route)

A route consists of a sequence of nodes  $\theta$  that starts and ends at the depot, where each node is associated with a (possibly empty) set of scheduled requests

### Definition (Budget of a route $\theta$ that started at instant $\tau$ )

The budget of a route  $\theta$  that started at instant  $\tau$ ,  $b(\tau, \theta)$ , is the slack time relative to the end of the service period (given by  $U$ )

## MDP Model: Decision Epochs and States

### Definition (Vehicle state)

State of vehicle  $k \in \{1, \dots, K\}$  is given by

$$V_k = \begin{cases} \emptyset & \text{if } k \text{ is idle (stationed at the depot)} \\ (\tau_k, \theta_k) & \text{if } k \text{ started to travel along route } \theta_k \text{ at instant } \tau_k \end{cases}$$

### Definition (State)

State  $S_0$  (initial state) represents all parameters of the problem

State  $S_t$ ,  $t \in \{1, \dots, T\}$ , is a  $(K + 1)$ -tuple  $S_t = (V_1, \dots, V_K, r_t)$ , where  $V_k$  are vehicle states, and  $r_t = (u_t, i_t, d_t)$  is the  $t$ -th element of  $\mathcal{D}(U)$



## Online Decisions: Scheduling Policies

At each state  $S_t$ ,  $t \in \{1, \dots, T\}$ , a scheduling policy prescribes:

- 1 **Acceptance** decision: accept or reject request  $r_t$
- 2 **Assignment** decision: if  $r_t$  is accepted, assign  $r_t$  to a vehicle  $k$
- 3 **Routing** decision: when  $r_t$  is assigned to vehicle  $k$ , define how route  $\theta_k$  is adjusted (in case  $V_k = (\tau_k, \theta_k)$ ) or initialized (in case  $V_k = \emptyset$ ) to accommodate  $r_t$



## Online Decisions: Routing Policies

Consider a request  $r = (u, i, d)$  and a vehicle state  $V = (\tau, \theta)$

### Definition (Cheapest Insertion (CI) Routing Policy $\rho_{\text{CI}}$ )

Routing policy  $\rho_{\text{CI}}$  inserts  $r$  into  $\theta$  by cheapest insertion, creating a new route  $\rho_{\text{CI}}(\theta, r)$

→ Polynomial time

### Definition (Reoptimization Routing Policy $\rho_{\text{R}}$ )

Routing policy  $\rho_{\text{R}}$  inserts  $r$  into  $\theta$ , creating a new route  $\theta' = \rho_{\text{R}}(\theta, r)$  such that the budget  $b(\tau, \theta')$  is maximized

→ NP-hard



## Online Decisions: State Transitions

- Let  $\mathcal{X}_t$  be the set of possible decisions when at state  $S_t$
- A policy  $\pi$  maps each possible state  $S_t$  to a decision  $X^\pi(S_t) \in \mathcal{X}_t$
- State transition:

$$S_{t+1} = S^M(S_t, x_t, r_{t+1})$$

- Decisions must be taken in “real-time” (i.e., in seconds, not minutes)



## MDP Model: Rewards and Objective

### Rewards:

$$R(S_t, x) = \begin{cases} 1 & \text{if } x \in \mathcal{X}_t \text{ is an 'accept' decision} \\ 0 & \text{otherwise} \end{cases}$$

### Objective:

$$\max_{\pi \in \Pi, y \in \mathcal{F}} \mathbb{E} \left[ \mathbb{E} \left[ \sum_{t=1}^T R(S_t, X^\pi(S_t)) \middle| S_1 \right] \right]$$

where  $\Pi$  is the set of feasible policies and  $\mathcal{F}$  is the set of feasible initial plans



## Online decisions:

- **Potential-based policy (PbP)**
- Simplified PbP (S-PbP)

## Offline decisions:

- Myopic plan
- Potential-based plan





## Potential of a State

Given a policy  $\pi$ , the potential of a state  $S_t$  is the expected reward-to-go:

$$\Phi_{\pi}(S_t) = \mathbb{E} \left[ \sum_{t'=t}^T R(S_{t'}, X^{\pi}(S_{t'})) \middle| S_t \right] \quad t \neq 0$$

Given an approximation function  $\hat{\Phi}_{\pi^*}(S_t) \approx \Phi_{\pi^*}(S_t)$  and a set of candidate decisions  $\tilde{\mathcal{X}}_t$ , we prescribe decision

$$X^{\pi}(S_t) = \arg \max_{x_t \in \tilde{\mathcal{X}}_t} R(S_t, x_t) + \mathbb{E} \left[ \hat{\Phi}_{\pi^*}(S_{t+1}) \middle| x_t \right]$$



# Towards an Approximation Model: Challenges

## Goal:

- Estimate the potential  $\Phi_{\pi^*}(S_t)$  of the optimal policy  $\pi^*$  for any given state  $S_t$

## What we have available:

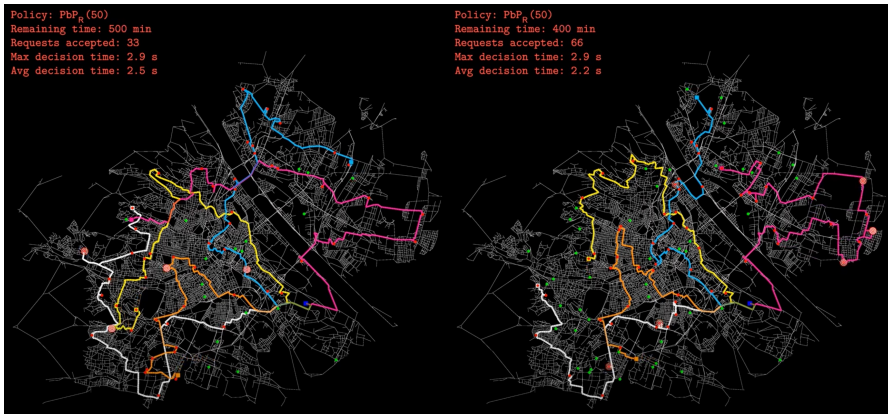
- We can sample trajectories  $\Omega = \{\omega_1, \dots, \omega_H\}$  from the spatiotemporal request distribution

## Two main challenges must be overcome:

- How to estimate the minimum budget required for serving a (stochastic) future request?
- How to handle the “competition” of vehicles to serve requests?



# Towards an Approximation Model: Challenges





## Towards an Approximation Model: Two Remarks

### Remark 1 (Late Depot Arrival)

*Since the goal is to maximize the number of accepted requests, drivers return to the depot at an instant near the end of the service period*

### Remark 2 (Request Order Preservation)

*The relative order of scheduled requests does not change (CI policy), or changes only slightly (reopt. policy), when a new request is assigned to a route*



## Modeling Insight: Effective Vehicle Speed

Remarks 1 and 2 suggest the concept of **effective speed**:

### Definition (Effective Vehicle Speed)

At instant  $u$ , the effective speed of a vehicle  $k$  with state  $V_k = (\tau_k, \theta_k)$  is the speed such that  $k$  arrives at the depot exactly at instant  $U$ , provided that  $V_k$  never changes

To predict the location of vehicle  $k$  at a future instant  $u' > u$ , we simulate route  $\theta_k$  under the effective speed for an amount of time  $u' - u$



## Modeling Insight: Assignment Costs

We now have a good prediction of where vehicle  $k$  will be at any future instant  $u'$ . Next, we determine:

$\mathcal{V}_{u_t}(V_k, u')$  (Predicted) set of nodes along  $\theta_k$  not yet traversed by vehicle  $V_k = (\tau_k, \theta_k)$  at instant  $u'$

Finally: the min cost (or budget consumption) when assigning a future request  $r' = (u', i', d')$  to vehicle  $k$  is approximated by:

$$c_{u_t}(V_k, r') = \min_{j \in \mathcal{V}_{u_t}(V_k, u')} t(j, i') + t(i', j) + d'$$



## Multiple-Knapsack Approximation of the Potential

Given  $H$  sample paths  $\Omega = \{\omega_1, \dots, \omega_H\}$  for the remaining horizon:

$$\mathbb{E} \left[ \hat{\Phi}_{\pi^*}(S_{t+1}) \middle| x_t \right] \approx \frac{1}{H} \sum_{\omega \in \Omega} \phi^{\omega}(S_{t+1} | x_t)$$

where

$$\begin{aligned} \phi^{\omega}(S_{t+1} | x_t) = \max \quad & \sum_{k \in \overline{K}} \sum_{r \in \omega} z_{kr} \\ \text{s.t.} \quad & \sum_{r \in \omega} c_{u_t}(V_k, r) z_{kr} \leq b(\tau_k, \theta_k) & k \in \overline{K} \\ & \sum_{k \in \overline{K}} z_{kr} \leq 1 & r \in \omega \\ & 0 \leq z_{kr} \leq 1 & k \in \overline{K}, r \in \omega \end{aligned}$$



## Potential-based Policy

**Given a routing policy  $\rho \in \{\rho_{\text{Cl}}, \rho_{\text{R}}\}$ . Upon arrival of  $r_t = (u_t, i_t, d_t)$ :**

- 1 Initialize a set of candidate decisions  $\tilde{\mathcal{X}}_t = \{x_t^-\}$  with the 'reject' decision  $x_t^-$  related to request  $r_t$
- 2 If it is feasible to serve  $r_t$  with vehicle  $k$  under the given routing policy, add the corresponding decision  $x_t^k$  to set  $\tilde{\mathcal{X}}_t$
- 3 Return decision

$$\arg \max_{x_t \in \tilde{\mathcal{X}}_t} R(S_t, x_t) + \frac{1}{H} \sum_{\omega \in \Omega} \phi^\omega(S_{t+1} | x_t)$$





## Potential-based Policy: Discussion

### Pros:

- Accuracy: MPE of  $\pm 2.5\%$  on most instances, from  $u = 0$
- Fast: requires the solution of  $H(K + 1)$  linear programs (LPs)
- **Zero** tunable parameters

### Cons:

- Not **so** fast: for very large  $K$  and very high request rate, it takes a while to solve all LPs
- The simplified PbP (S-PbP) trades-off accuracy by efficiency

### What sort of policy is that?

- <http://tinyurl.com/Powelllookaheadpolicies>
- Stochastic lookahead policy with sampling, stage aggregation, latent variables and policy approximation



## Benchmark Policies

- Greedy policies:  $GP_{Cl}$  and  $GP_R$ 
  - Accept all feasible requests; with and without reoptimization
- Rollout:  $R_{Cl}-GP_{Cl}(H)$  and  $R_R-GP_{Cl}(H)$ 
  - $H$  sample paths,  $GP_{Cl}$  as base policy
- Policy function approximation:  $PFA_{Cl}$  and  $PFA_R$ 
  - Single parameter (trained offline) trades off immediate reward and reward-to-go
- Rollout:  $R_{Cl}-PFA_{Cl}(H)$  and  $R_R-PFA_{Cl}(H)$ 
  - $H$  sample paths,  $PFA_{Cl}$  as base policy



# Computational Study

- Real street network of Vienna (16,080 nodes and 36,424 arcs)
- Service period: 10 hours
- Number of vehicles:  $K \in \{2, 3, 5, 6, 10, 12, 20\}$
- Request rate (per minute):  $\Lambda \in \{0.2, 0.4, 0.8, 1.5\}$ 
  - Smallest instances larger than most instances from previous works
- Degree of dynamism:  $\eta \in \{75\%, 85\%, 90\%, 95\%\}$
- Three spatiotemporal request distributions:
  - Uniform, time-independent (UTI)
  - Clustered, time-independent (CTI)
  - Clustered, time-dependent (CTD)



# Instances, Policies and Offline Planners Simulated

Instance parameters					Algorithms		Simulations
$\Lambda$	$\eta$	$K$	Dist.	Scen.	Offline	Online	
0.2	0.75	2, 3	UTI	5	MY PB	$GP_{Cl}, GP_R, R_{Cl}-GP_{Cl}(10, 25, 50, 100),$	12,270
			CTI			$R_R-GP_{Cl}(10, 25, 50, 100), PFA_{Cl}, PFA_R,$	
			CTD			$R_R-PFA_{Cl}(25, 50, 100)^*, S-PbP, PbP$	
0.4	0.85	3, 5	UTI	5	MY PB	$GP_{Cl}, GP_R, R_{Cl}-GP_{Cl}(10, 25, 50, 100),$	12,270
			CTI			$R_R-GP_{Cl}(10, 25, 50, 100), PFA_{Cl}, PFA_R,$	
			CTD			$R_R-PFA_{Cl}(25, 50, 100)^*, S-PbP, PbP$	
0.8	0.90	6, 12	UTI	5	PB	$GP_{Cl}, GP_R, PFA_{Cl}, PFA_R,$	3,210
			CTI			$R_{Cl}-GP_{Cl}(10), R_R-GP_{Cl}(10),$	
			CTD			$R_R-PFA_{Cl}(10), S-PbP, PbP$	
1.5	0.95	10, 20	UTI	1**	PB	$GP_{Cl}, GP_R, PFA_{Cl}, PFA_R,$	214
						$R_{Cl}-GP_{Cl}(10), R_R-GP_{Cl}(10),$	
						$R_R-PFA_{Cl}(10), S-PbP, PbP$	




## Potential-based vs Myopic Plans

Average number of accepted requests:

Policy	Myopic Plan	Potential-based Plan	Diff
GP <sub>Cl</sub>	54.1	61.2	13.0%
GP <sub>R</sub>	54.9	62.7	14.1%
R <sub>Cl</sub> -GP <sub>Cl</sub> (10)	63.0	66.2	5.1%
R <sub>Cl</sub> -GP <sub>Cl</sub> (25)	65.8	69.4	5.4%
R <sub>Cl</sub> -GP <sub>Cl</sub> (50)	67.3	71.1	5.6%
R <sub>Cl</sub> -GP <sub>Cl</sub> (100)	68.2	72.1	5.8%
R <sub>R</sub> -GP <sub>Cl</sub> (10)	63.4	67.2	6.1%
R <sub>R</sub> -GP <sub>Cl</sub> (25)	66.2	70.7	6.7%
R <sub>R</sub> -GP <sub>Cl</sub> (50)	67.7	72.2	6.8%
R <sub>R</sub> -GP <sub>Cl</sub> (100)	68.5	73.3	7.0%
S-PbP(50)	72.0	75.2	4.4%
<b>PbP(50)</b>	<b>72.9</b>	<b>77.2</b>	<b>5.9%</b>
PFA <sub>Cl</sub>	61.3	71.8	17.1%
PFA <sub>R</sub>	62.1	72.6	16.9%
Avg	66.4	72.4	9.0%




## Policy Comparison ( $\Lambda \in \{0.2, 0.4\}$ )









## Policy Comparison ( $\Lambda \in \{0.8, 1.5\}$ )





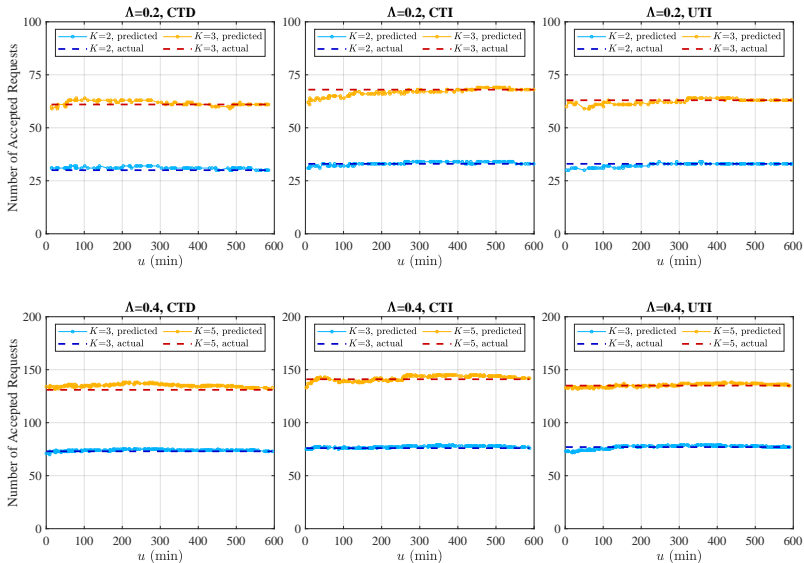
# Acceptance Profiles







# Multiple-Knapsack Potential Approximations





# Conclusions

- Main contributions and takeaways:
  - Expected reward-to-go can be accurately and efficiently approximated by knapsack models
  - Accurate potential approximations enable high-performing scheduling policies, which outperform classical ADP methods traditionally used for DVRPs such as rollout algorithms and PFA
  - Coverage of the service area is more important than budget alone
- Possible extensions and future research:
  - Vehicle capacity, time windows, (self-imposed) time window assignments
  - Pickup and delivery
  - Reassignment of requests among vehicles (more complex)



## Acknowledgements

This project has received funding from the European Union's Horizon 2020 research and innovation programme under the Marie Skłodowska-Curie grant agreement No 754462

Computational experiments were performed on the Dutch national e-infrastructure with the support of SURF Cooperative

Geographical data for Vienna are copyrighted to OpenStreetMap contributors and available at <http://openstreetmap.org>

## Questions & Discussions